

# Data Fragmentation

---

**Due** Mar 23 by 11:59pm    **Points** 20    **Submitting** a text entry box or a file upload  
**Available** Mar 9 at 12am - Mar 23 at 11:59pm 15 days

---

This assignment was locked Mar 23 at 11:59pm.

The required task is to simulate data partitioning approaches on-top of an open source relational database management system (i.e., PostgreSQL). Each student must generate a set of Python functions that load the input data into a relational table, partition the table using different horizontal fragmentation approaches, and insert new tuples into the right fragment. A detailed explanation about round-robin partitioning and range partitioning can be found here:

[https://www.ibm.com/support/knowledgecenter/en/SSZJPZ\\_11.7.0/com.ibm.swg.im.iis.ds.parjob.dev.doc/topics/partitioning.html](https://www.ibm.com/support/knowledgecenter/en/SSZJPZ_11.7.0/com.ibm.swg.im.iis.ds.parjob.dev.doc/topics/partitioning.html)

([https://www.ibm.com/support/knowledgecenter/en/SSZJPZ\\_11.7.0/com.ibm.swg.im.iis.ds.parjob.dev.doc/topics/partitioning.html](https://www.ibm.com/support/knowledgecenter/en/SSZJPZ_11.7.0/com.ibm.swg.im.iis.ds.parjob.dev.doc/topics/partitioning.html))

**Input Data.** The input data is a Movie Rating data set collected from the MovieLens web site (<http://movielens.org>). The raw data is available in the file ratings.dat.

The rating.dat file contains 10 million ratings and 100,000 tag applications applied to 10,000 movies by 72,000 users. Each line of this file represents one rating of one movie by one user, and has the following format:

UserID::MovieID::Rating::Timestamp

Ratings are made on a 5-star scale, with half-star increments. Timestamps represent seconds since midnight Coordinated Universal Time (UTC) of January 1, 1970. A sample of the file contents is given below:

1::122::5::838985046

1::185::5::838983525

1::231::5::838983392

**Required Task.** Below are the steps you need to follow to fulfill this assignment:

1. Install PostgreSQL.
2. Download rating.dat file from the MovieLens website, <http://files.grouplens.org/datasets/movielens/ml-10m.zip>

You can use partial data for testing.

3. Implement a Python function Load Ratings() that takes a file system path that contains the rating.dat file as input. Load Ratings() then load the rating.dat content into a table (saved in PostgreSQL) named Ratings that has the following schema

UserID(int) – MovieID(int) – Rating(float)

4. Implement a Python function Range\_Partition() that takes as input: (1) the Ratings table stored in PostgreSQL and (2) an integer value N; that represents the number of partitions. Range\_Partition() then generates N horizontal fragments of the Ratings table and store them in PostgreSQL. The algorithm should partition the ratings table based on N uniform ranges of the Rating attribute.

5. Implement a Python function RoundRobin\_Partition() that takes as input: (1) the Ratings table stored in PostgreSQL and (2) an integer value N; that represents the number of partitions. The function then generates N horizontal fragments of the Ratings table and stores them in PostgreSQL. The algorithm should partition the ratings table using the round robin partitioning approach (explained in class).

6. Implement a Python function RoundRobin\_Insert() that takes as input: (1) Ratings table stored in PostgreSQL, (2) UserID, (3) ItemID, (4) Rating. RoundRobin\_Insert() then inserts a new tuple to the Ratings table and the right fragment based on the round robin approach.

7. Implement a Python function Range\_Insert() that takes as input: (1) Ratings table stored in PostgreSQL (2) UserID, (3) ItemID, (4) Rating. Range\_Insert() then inserts a new tuple to the Ratings table and the correct fragment (of the partitioned ratings table) based upon the Rating value.

8. Implement function Delete\_Partitions() for your testing convenience. It will not be graded.

### Frequently Asked Questions:

- Partition numbers start from 0, if there are 3 partitions then range\_part0, range\_part1, range\_part2 are partition table names for range partitions and similar numbering should be done for round robin partitions.
- Do not change partition table names prefix given in assignment\_tester.py
- Do not hard code input filename.
- Do not hard code database name.
- Table schema should be equivalent to what has been described in point 3.
- Use Python 2.7.x version.

### Partitioning Questions:

The number of partitions here refer to the number of tables to be created. For rating values in [0, 0.5, 1, 1.5, 2, 2.5, 3, 3.5, 4, 4.5, 5]

Case N = 1,

One table containing all the values.

Case N = 2, Two tables,

Partition 0 has values [0,2.5]

Partition 1 has values (2.5,5]

Case N = 3, Three tables,

Partition 0 has values [0, 1.67]

Partition 1 has values (1.67, 3.34]

Partition 2 has values (3.34, 5]

Uniform ranges means a region is divided uniformly, I hope the example gives a clear picture.

### Assignment Tips!

1. Do not use global variables in your implementation. Meta-data table in the database is allowed.
2. You are not allowed to modify the data file on disk.
3. Two insert functions can be called many times at any time. They are designed for maintaining the tables in the database when insertions happen.

### Submission

Only submit the Interface.py file. Do not change the file name. Do not put it into a folder or upload a zip.

We provide a virtual machine that has the testing environment and an installed PostgreSQL.

OS username: user

OS password: user

Postgres username: postgres

Postgres password: 1234

You can download it and use any VM software such as **VirtualBox** (<https://www.virtualbox.org/>) to import it:

[https://drive.google.com/drive/folders/1\\_efzBidl7waVEjVxJw2OJ3kv2PZO-dNm?usp=sharing](https://drive.google.com/drive/folders/1_efzBidl7waVEjVxJw2OJ3kv2PZO-dNm?usp=sharing)

You will use the following files within your assignment. These files are used to test your Interface.py

1. tester.py: Test your interface.py using this tester. Run it using "python tester.py".

- [tester.py](https://canvas.asu.edu/courses/75440/files/26129836/download?download_frd=1)  (https://canvas.asu.edu/courses/75440/files/26129836/download?download\_frd=1)


2. testHelper.py: put this one together with tester.py

- [testHelper.py](https://canvas.asu.edu/courses/75440/files/30185251/download?download_frd=1)  (https://canvas.asu.edu/courses/75440/files/30185251/download?download\_frd=1)

3. test\_data.txt: some test data

- [test\\_data.txt](https://canvas.asu.edu/courses/75440/files/30185277/download?download_frd=1)  (https://canvas.asu.edu/courses/75440/files/30185277/download?download\_frd=1)

4. Interface.py: Implement the interface in Interface.py

- [Interface.py](https://canvas.asu.edu/courses/75440/files/30185335/download?download_frd=1)  (https://canvas.asu.edu/courses/75440/files/30185335/download?download\_frd=1)